

✓ Alexandre Delode

Research Positions in Machine Learning and Computer Vision

June 19, 2024

✓ Task 1

Import the required libraries

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, UpSampling
3 from tensorflow.keras.models import Model
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7
```

Preprocess the dataset

```
1 import tensorflow as tf
2 import tensorflow_datasets as tfds
3 import numpy as np
4
5 # Define image size
6 image_size = 64
7
8 # Load the dataset
9 dataset, info = tfds.load('tf_flowers', with_info=True, as_supervised=True)
10 train_dataset = dataset['train'].shuffle(1000).take(80) # Use 80 images for
11 test_dataset = dataset['train'].shuffle(1000).skip(80).take(20) # Use 20 i
12
13 # Preprocess function
14 def preprocess(image, label):
15     # Resize and normalize images
16     image = tf.image.resize(image, (image_size, image_size))
17     image = tf.image.rgb_to_grayscale(image) # Convert to grayscale
18     image = tf.cast(image, tf.float32) / 255.0 # Normalize pixel values to
19     return image, label
20
21 # Apply preprocess function and batch the datasets
22 batch_size = 32
23 train_dataset = train_dataset.map(preprocess).batch(batch_size)
24 test_dataset = test_dataset.map(preprocess).batch(batch_size)
```

```

25
26 # Convert train dataset to numpy arrays
27 train_images = []
28 train_labels = []
29 for images, labels in tfds.as_numpy(train_dataset):
30     train_images.append(images)
31     train_labels.append(labels)
32
33 train_images = np.concatenate(train_images)
34 train_labels = np.concatenate(train_labels)
35
36 # Convert test dataset to numpy arrays
37 test_images = []
38 test_labels = []
39 for images, labels in tfds.as_numpy(test_dataset):
40     test_images.append(images)
41     test_labels.append(labels)
42
43 test_images = np.concatenate(test_images)
44 test_labels = np.concatenate(test_labels)
45
46 # Print shapes to verify
47 print('Train images shape:', train_images.shape)
48 print('Train labels shape:', train_labels.shape)
49 print('Test images shape:', test_images.shape)
50 print('Test labels shape:', test_labels.shape)

```

```

➡ Downloading and preparing dataset 218.21 MiB (download: 218.21 MiB, generated: 2
Dl Completed...: 0%|          | 0/5 [00:00<?, ? file/s]
Dataset tf_flowers downloaded and prepared to /root/tensorflow_datasets/tf_flowe
Train images shape: (80, 64, 64, 1)
Train labels shape: (80,)
Test images shape: (20, 64, 64, 1)
Test labels shape: (20,)

```

Add uniform noise for the training set.

```

1 def add_uniform_noise(images, noise_factor=0.2):
2     noisy_images = images + noise_factor * np.random.uniform(low=-1.0, high=1.0, size=ir
3     noisy_images = np.clip(noisy_images, 0., 1.)
4     return noisy_images
5
6 noisy_train_images = add_uniform_noise(train_images)
7

```

Create one model that includes the encoder and the decoder.

```

1 from tensorflow.keras import layers, Model, losses, callbacks
2
3 # Define image size
4 image_size = 64
5
6 class Denoise(Model):
7     def __init__(self):
8         super(Denoise, self).__init__()
9         self.encoder = tf.keras.Sequential([
10             layers.Input(shape=(image_size, image_size, 1)),
11             layers.Conv2D(32, (3, 3), activation='relu', padding='same', strides=2),
12             layers.Conv2D(18, (3, 3), activation='relu', padding='same', strides=2)])
13
14         self.decoder = tf.keras.Sequential([
15             layers.Conv2DTranspose(18, kernel_size=3, strides=2, activation='relu', padding='same'),
16             layers.Conv2DTranspose(32, kernel_size=3, strides=2, activation='relu', padding='same'),
17             layers.Conv2D(1, kernel_size=(3, 3), activation='sigmoid', padding='same')])
18
19     def call(self, x):
20         encoded = self.encoder(x)
21         decoded = self.decoder(encoded)
22         return decoded
23
24 # Build and summarize the autoencoder
25 autoencoder = Denoise()
26 input_img = tf.keras.Input(shape=(image_size, image_size, 1))
27 output_img = autoencoder(input_img)
28 autoencoder_model = Model(input_img, output_img)
29 autoencoder_model.compile(optimizer='adam', loss=losses.MeanSquaredError())
30 autoencoder_model.summary()

```

➡ Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 64, 64, 1)]	0
denoise (Denoise)	(None, 64, 64, 1)	13961
=====		
Total params: 13961 (54.54 KB)		
Trainable params: 13961 (54.54 KB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

The structure has been taken from the tensorflow tutorial. It has been adapted to fit a 64 x 64 image. The hyperparameters could have been optimised with optuna.

```

1
2 # Define early stopping callback
3 early_stopping = callbacks.EarlyStopping(
4     monitor='val_loss',      # Monitor validation loss
5     patience=10,             # Number of epochs with no improvement after which training
6     restore_best_weights=True # Restore the weights from the epoch with the best value
7 )
8
9 # Train the autoencoder with a history callback
10 history = autoencoder_model.fit(noisy_train_images, train_images,
11                                 epochs=500,
12                                 batch_size=16,
13                                 shuffle=True,
14                                 validation_split=0.1,
15                                 callbacks=[early_stopping])
16

```



```

Epoch 1/500
5/5 [=====] - 2s 139ms/step - loss: 0.0758 - val_loss:
Epoch 2/500
5/5 [=====] - 0s 89ms/step - loss: 0.0736 - val_loss: 0
Epoch 3/500
5/5 [=====] - 0s 95ms/step - loss: 0.0696 - val_loss: 0
Epoch 4/500
5/5 [=====] - 1s 112ms/step - loss: 0.0658 - val_loss:
Epoch 5/500
5/5 [=====] - 1s 150ms/step - loss: 0.0603 - val_loss:
Epoch 6/500
5/5 [=====] - 1s 144ms/step - loss: 0.0534 - val_loss:
Epoch 7/500
5/5 [=====] - 1s 158ms/step - loss: 0.0455 - val_loss:
Epoch 8/500
5/5 [=====] - 1s 142ms/step - loss: 0.0375 - val_loss:
Epoch 9/500
5/5 [=====] - 1s 138ms/step - loss: 0.0306 - val_loss:
Epoch 10/500
5/5 [=====] - 1s 102ms/step - loss: 0.0261 - val_loss:
Epoch 11/500
5/5 [=====] - 0s 88ms/step - loss: 0.0231 - val_loss: 0
Epoch 12/500
5/5 [=====] - 0s 101ms/step - loss: 0.0209 - val_loss:
Epoch 13/500
5/5 [=====] - 0s 92ms/step - loss: 0.0188 - val_loss: 0
Epoch 14/500
5/5 [=====] - 0s 96ms/step - loss: 0.0177 - val_loss: 0
Epoch 15/500
5/5 [=====] - 0s 92ms/step - loss: 0.0168 - val_loss: 0
Epoch 16/500
5/5 [=====] - 0s 93ms/step - loss: 0.0161 - val_loss: 0
Epoch 17/500
5/5 [=====] - 0s 92ms/step - loss: 0.0158 - val_loss: 0
Epoch 18/500
5/5 [=====] - 0s 86ms/step - loss: 0.0153 - val_loss: 0
Epoch 19/500
5/5 [=====] - 0s 95ms/step - loss: 0.0150 - val_loss: 0

```

```

Epoch 20/500
5/5 [=====] - 0s 91ms/step - loss: 0.0147 - val_loss: 0
Epoch 21/500
5/5 [=====] - 0s 97ms/step - loss: 0.0147 - val_loss: 0
Epoch 22/500
5/5 [=====] - 0s 88ms/step - loss: 0.0143 - val_loss: 0
Epoch 23/500
5/5 [=====] - 0s 96ms/step - loss: 0.0139 - val_loss: 0
Epoch 24/500
5/5 [=====] - 0s 97ms/step - loss: 0.0137 - val_loss: 0
Epoch 25/500
5/5 [=====] - 0s 100ms/step - loss: 0.0136 - val_loss:
Epoch 26/500
5/5 [=====] - 0s 89ms/step - loss: 0.0135 - val_loss: 0
Epoch 27/500
5/5 [=====] - 0s 95ms/step - loss: 0.0132 - val_loss: 0
Epoch 28/500
5/5 [=====] - 0s 94ms/step - loss: 0.0131 - val_loss: 0
Epoch 29/500
5/5 [=====] - 0s 88ms/step - loss: 0.0130 - val_loss: 0


```

The test data are kept in a "vault" and not used for the training.

```

1 # Plot training & validation loss values
2 plt.figure(figsize=(10, 5))
3 plt.plot(history.history['loss'], label='Train Loss')
4 plt.plot(history.history['val_loss'], label='Validation Loss')
5 plt.title('Model loss')
6 plt.xlabel('Epoch')
7 plt.ylabel('Loss')
8 plt.legend(loc='upper right')
9 plt.show()

```



There is a plateau. The training is stopped before overfitting.

```
1 def add_gaussian_noise(images, noise_factor=0.2):
2     noisy_images = images + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=images.shape)
3     noisy_images = np.clip(noisy_images, 0., 1.)
4     return noisy_images
5
6 noisy_test_images = add_gaussian_noise(test_images)
```

Gaussian noise is added to the test set.

```
1 # Predict denoised images
2 denoised_images = autoencoder_model.predict(noisy_test_images)
```

↩ 1/1 [=====] - 0s 336ms/step

```
1 from skimage.metrics import mean_squared_error, structural_similarity
2
3 # Calculate MSE for each image in the test set
4 mse_values = [mean_squared_error(original, reconstructed) for original, reconstructed in zip(test_images, denoised_images)]
5
6 # Calculate SSIM for each image in the test set
7 ssim_values = [structural_similarity(original.reshape(64, 64), reconstructed.reshape(64, 64)) for original, reconstructed in zip(test_images, denoised_images)]
8
9 # Calculate MSE for each noisy image in the test set
10 mse_values_noisy = [mean_squared_error(original, reconstructed) for original, reconstructed in zip(noisy_test_images, denoised_images)]
11
12 # Calculate SSIM for each noisy image in the test set
13 ssim_values_noisy = [structural_similarity(original.reshape(64, 64), reconstructed.reshape(64, 64)) for original, reconstructed in zip(noisy_test_images, denoised_images)]
14
15 # Print average MSE and SSIM
16 print("Average MSE noisy:", np.mean(mse_values_noisy))
17 print("Average SSIM noisy:", np.mean(ssim_values_noisy))
18
19 # Print average MSE and SSIM
20 print("Average MSE denoised:", np.mean(mse_values))
21 print("Average SSIM denoised:", np.mean(ssim_values))
```

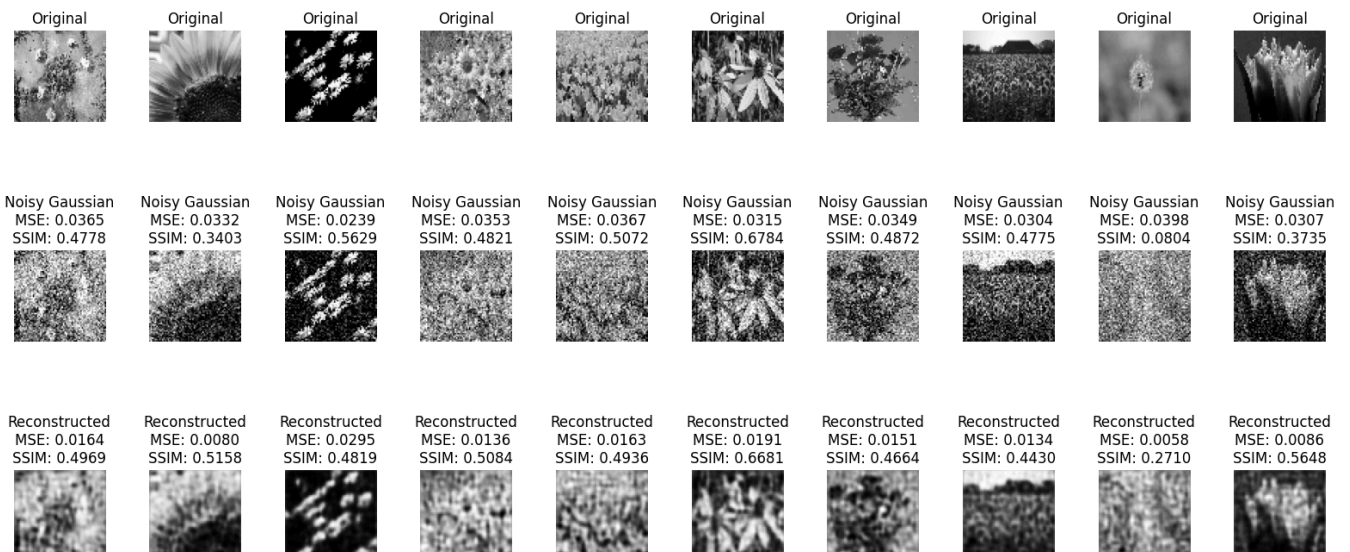
↩ Average MSE noisy: 0.03308295542824337
Average SSIM noisy: 0.398677768543888
Average MSE denoised: 0.012956395754919631
Average SSIM denoised: 0.4986860144437931

Even if the model is limited in parameters, the training is limited by Colab capabilities and the dataset is relatively small, an improvement can be seen from the MSE and SSIM metrics with MSE divided by more than two and SSIM augmented by 50%.

```

1 # Visualize the results
2 n = 10
3 plt.figure(figsize=(20, 8))
4 for i in range(n):
5     # Display original
6     ax = plt.subplot(5, n, i + 1)
7     plt.imshow(test_images[i].reshape(image_size, image_size), cmap='gray')
8     plt.title("Original")
9     plt.axis('off')
10
11     # Display noisy
12     ax = plt.subplot(5, n, i + 1 + 2*n)
13     plt.imshow(noisy_test_images[i].reshape(image_size, image_size), cmap='gray')
14     plt.title(f"Noisy Gaussian\nMSE: {mse_values_noisy[i]:.4f}\nSSIM: {ssim_values_noisy[i]:.4f}")
15     plt.axis('off')
16
17     # Display denoised
18     ax = plt.subplot(5, n, i + 1 + 4*n)
19     plt.imshow(denoised_images[i].reshape(image_size, image_size), cmap='gray')
20     plt.title(f"Reconstructed\nMSE: {mse_values[i]:.4f}\nSSIM: {ssim_values[i]:.4f}")
21     plt.axis('off')
22
23 plt.show()

```



We can visually see the improvement from the noisy image to the denoised image compared to the original one.

✓ Task 2 - First Part

Import the required library

```

1 import tensorflow as tf
2 from tensorflow.keras import layers, Model, losses
3 from tensorflow.keras.datasets import cifar10
4 import numpy as np
5 import matplotlib.pyplot as plt

```

Import CIFAR dataset and create train/test sets.

```

1 # Load CIFAR-10 dataset
2 (x_train, _), (_, _) = cifar10.load_data()
3
4 # Take 100 images for the task
5 images = x_train[:100]
6 images_test = x_train[100:200]

```

```

1 # Visualize the dataset
2 n = 10
3 plt.figure(figsize=(20, 5))
4 for i in range(n):
5     # Display original
6     ax = plt.subplot(5, n, i + 1)
7     plt.imshow(images[i])
8     plt.title("Original")
9     plt.axis('off')

```



"To make things more interesting, we also augment the images before feeding them to R.": So the data augmentation must be coded as a function.

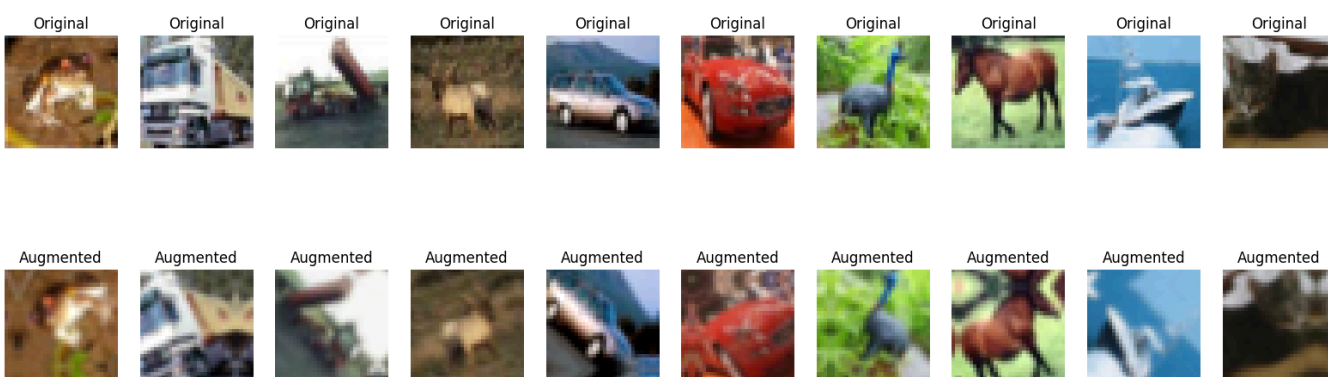
```

1 def data_augmentation(images):
2     # Function to perform data augmentation
3     data_aug = tf.keras.Sequential([
4         layers.experimental.preprocessing.RandomFlip("horizontal"),
5         layers.experimental.preprocessing.RandomRotation(0.2),
6         layers.experimental.preprocessing.RandomTranslation(0.2, 0.2),
7         layers.experimental.preprocessing.RandomContrast(0.2),
8         layers.GaussianNoise(0.1)
9     ])
10    # Apply augmentation to each image
11    augmented_images = np.array([tf.cast(data_aug(image), tf.float32) / 255.0 for image
12                                     in images])
13    return augmented_images
14
15 # Augment images
16 augmented_images = data_augmentation(images)
17

```


Visualize the augmented data set.

```
1 # Visualize the results
2 n = 10
3 plt.figure(figsize=(20, 10))
4 for i in range(n):
5     # Display original
6     ax = plt.subplot(3, n, i + 1)
7     plt.imshow(images[i])
8     plt.title("Original")
9     plt.axis('off')
10
11     # Display noisy
12     ax = plt.subplot(3, n, i + 1 + n)
13     plt.imshow(augmented_images[i])
14     plt.title("Augmented")
15     plt.axis('off')
```



**Due to a limited amount of time with the only use of Colab, I had to work in a different file that I can't merge in order to keep the compilation. **

Alexandre Delode

Research Positions in Machine Learning and Computer Vision

June 19, 2024

Task 2 - Second part

Import the required libraries

```
In [ ]: import tensorflow as tf
        from tensorflow.keras import layers, Model, losses
        from tensorflow.keras.datasets import cifar10
        import numpy as np
        import matplotlib.pyplot as plt
```

Create train/test sets

```
In [ ]: # Load CIFAR-10 dataset
        (x_train, _), (_, _) = cifar10.load_data()

        # Take 100 images for the task
        images = x_train[:100]
        images_test = x_train[100:200]
```

The data augmentation , see first file for some exemples

```
In [ ]: def data_augmentation(images):
        # Function to perform data augmentation
        data_aug = tf.keras.Sequential([
            layers.experimental.preprocessing.RandomFlip("horizontal"),
            layers.experimental.preprocessing.RandomRotation(0.2),
            layers.experimental.preprocessing.RandomTranslation(0.2, 0.2),
            layers.experimental.preprocessing.RandomContrast(0.2),
            layers.GaussianNoise(0.1)
        ])
        # Apply augmentation to each image
        augmented_images = np.array([tf.cast(data_aug(image), tf.float32) / 255.0 for image in images])

        return augmented_images

        # Augment images
        augmented_images = data_augmentation(images)
```

Some basic values about VGG16 for the feature extraction.

```
In [ ]: import tensorflow as tf
        from tensorflow.keras import layers, models, applications
        import numpy as np
```

```

input_shape = (32, 32, 3)
base_model = applications.VGG16(weights='imagenet', include_top=False, input_shape=input_shape)
features1 = base_model(tf.expand_dims(images[0], axis=0))
print(features1.shape)
print(tf.reduce_min(features1))
print(tf.reduce_max(features1))
features2 = base_model(tf.expand_dims(images[0], axis=0))
print(tf.reduce_min(features2))
print(tf.reduce_max(features2))

```

```

(1, 1, 1, 512)
tf.Tensor(0.0, shape=(), dtype=float32)
tf.Tensor(171.37624, shape=(), dtype=float32)
tf.Tensor(0.0, shape=(), dtype=float32)
tf.Tensor(171.37624, shape=(), dtype=float32)

```

The sender has the task to order the images. The features are extracted by VGG16. This ordering must be trainable in order to optimize the objective. The solution of matrix multiplication with a learnable vector has been chosen. For the backpropagation process, the order choice must be soft. It is done by a sigmoid function with a scaling factor.

```

In [ ]: class SenderModel(tf.keras.Model):
        def __init__(self):
            super(SenderModel, self).__init__()
            self.input1=layers.InputLayer(input_shape=(32,32,3))
            self.input2=layers.InputLayer(input_shape=(32,32,3))
            self.orderim = tf.Variable(tf.random.uniform(shape=(512,)), minval=-1.0, maxval=1.0)
            self.base_model = applications.VGG16(weights='imagenet', include_top=False, input_shape=input_shape)
            for layer in self.base_model.layers:
                layer.trainable = False

        def call(self, x1, x2):
            x1 = self.input1(x1)
            x2 = self.input2(x2)

            features1 = self.base_model(x1)
            features2 = self.base_model(x2)

            # Compute scalar product (dot product)
            score1 = tf.reduce_sum(self.orderim * features1, axis=[1, 2, 3], keepdims=True)
            score2 = tf.reduce_sum(self.orderim * features2, axis=[1, 2, 3], keepdims=True)

            # Differentiable comparison
            comparison = tf.keras.activations.sigmoid((score1 - score2)*10)
            return comparison

```

```

In [ ]: # Build and summarize the sender
Sender = SenderModel()
input_sender_img1 = tf.keras.Input(shape=(32,32,3))
input_sender_img2 = tf.keras.Input(shape=(32,32,3))
output_sender_ord = Sender(input_sender_img1,input_sender_img2)
sender_model = Model((input_sender_img1,input_sender_img2), output_sender_ord)
sender_model.compile(optimizer='adam', loss=losses.MeanSquaredError())
sender_model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
=====			

input_6 (InputLayer)	[(None, 32, 32, 3)]	0	[]
input_7 (InputLayer)	[(None, 32, 32, 3)]	0	[]
sender_model (SenderModel)	(None, 1, 1, 1)	1471520 0	['input_6[0][0]', 'input_7[0][0]']

```
=====
=====
```

```
Total params: 14715200 (56.13 MB)
Trainable params: 512 (2.00 KB)
Non-trainable params: 14714688 (56.13 MB)
```

Some tests to check the almost 0-1 output.

```
In [ ]: sender_model((tf.expand_dims(images[0], axis=0),tf.expand_dims(images[1], axis=0)))
```

```
Out[ ]: <tf.Tensor: shape=(1, 1, 1, 1), dtype=float32, numpy=array([[[[0.]]]], dtype=float32)>
```

```
In [ ]: sender_model((tf.expand_dims(images[1], axis=0),tf.expand_dims(images[0], axis=0)))
```

```
Out[ ]: <tf.Tensor: shape=(1, 1, 1, 1), dtype=float32, numpy=array([[[[1.]]]], dtype=float32)>
```

```
In [ ]: sender_model((tf.expand_dims(images[1], axis=0),tf.expand_dims(images[98], axis=0)))
```

```
Out[ ]: <tf.Tensor: shape=(1, 1, 1, 1), dtype=float32, numpy=array([[[[0.]]]], dtype=float32)>
```

The embedding is done like the assignment description with two trainable vectors. As the output of the sender is almost 0-1, the output of the embedding model is almost xi1 or xi2.

```
In [ ]: embed_dim=512

class EmbeddingModel(tf.keras.Model):
    def __init__(self):
        super(EmbeddingModel, self).__init__()
        self.xi_0 = tf.Variable(tf.random.uniform(shape=(embed_dim,)), minval=-1.0, maxval=1.0)
        self.xi_1 = tf.Variable(tf.random.uniform(shape=(embed_dim,)), minval=-1.0, maxval=1.0)
    def call(self, x):
        x = tf.reshape(x, [-1])
        y=(1-x)*self.xi_0 + x*self.xi_1

        return y

# Build and summarize the embedder
Embedding_1 = EmbeddingModel()
input_emb_order = tf.keras.Input(shape=(1))
output_emb_vect = Embedding_1(input_emb_order)
emb_model = Model(input_emb_order, output_emb_vect)
emb_model.compile(optimizer='adam', loss=losses.MeanSquaredError())
emb_model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_8 (InputLayer)	[(None, 1)]	0
embedding_model (Embedding (512,))		1024

Model)

```
=====
Total params: 1024 (4.00 KB)
Trainable params: 1024 (4.00 KB)
Non-trainable params: 0 (0.00 Byte)
```

The receiver is a neural network classifier. The inputs are the concatenation of the features of the 2 images and the vector xi1 or xi2. The sigmoid output is a probability of the transposition image.

```
In [ ]: class ReceiverModel(tf.keras.Model):
        def __init__(self):
            super(ReceiverModel, self).__init__()
            self.input1=layers.InputLayer(input_shape=(32,32,3))
            self.input2=layers.InputLayer(input_shape=(32,32,3))
            self.reshape = layers.Reshape((1, 1, embed_dim)) # Reshape layer to match CNN
            # Initialize VGG16
            self.base_model = applications.VGG16(weights='imagenet', include_top=False, i
            for layer in self.base_model.layers:
                layer.trainable = False
            self.normalize = layers.LayerNormalization(axis=-1)
            self.concat = layers.Concatenate()
            self.dense1=layers.Dense(256, activation='relu')
            self.drop=layers.Dropout(0.2)
            self.dense2=layers.Dense(128, activation='relu') # Output scalar value center
            self.dense3=layers.Dense(1, name="outputs", activation='sigmoid') # Output sc

        def call(self, x1 , x2, key):

            x1 = self.input1(x1)
            x2 = self.input2(x2)

            features1 = self.base_model(x1)
            features2 = self.base_model(x2)

            # Normalize features
            features1 = self.normalize(features1)
            features2 = self.normalize(features2)

            key = tf.reshape(key, (1, 1, embed_dim))
            key = self.reshape(key)
            combined = self.concat([features1, features2, key])
            combined = self.dense1(combined)
            combined = self.drop(combined)
            combined = self.dense2(combined)
            combined = self.dense3(combined)

            return combined

        # Build and summarize the receiver
        Receiver = ReceiverModel()
        input_rec_imgaug1 = tf.keras.Input(shape=(32,32,3))
        input_rec_imgaug2 = tf.keras.Input(shape=(32,32,3))
        input_rec_key = tf.keras.Input(shape=( embed_dim))
        output_rec_vect = Receiver(input_rec_imgaug1,input_rec_imgaug2,input_rec_key)
        rec_model = Model((input_rec_imgaug1,input_rec_imgaug2,input_rec_key), output_rec_vec
        rec_model.compile(optimizer='adam', loss=losses.MeanSquaredError())
        rec_model.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_18 (InputLayer)	[(None, 32, 32, 3)]	0	[]
input_19 (InputLayer)	[(None, 32, 32, 3)]	0	[]
input_20 (InputLayer)	[(None, 512)]	0	[]
receiver_model_1 (Receiver Model)	(1, 1, 1, 1)	15142209	['input_18[0][0]', 'input_19[0][0]', 'input_20[0][0]']
=====			
=====			
Total params: 15142209 (57.76 MB)			
Trainable params: 427521 (1.63 MB)			
Non-trainable params: 14714688 (56.13 MB)			

Some methods for partial saving due to Colab disconnection.

```
In [ ]: import json
cumaccuracies=[]
cumlosses=[]
def save_to_json(lists_to_write, file_path):
    # Get the values of the lists from the global namespace
    data = {listname: globals()[listname] for listname in lists_to_write}

    # Save the data to a JSON file
    with open(file_path, 'w') as file:
        json.dump(data, file, indent=2)

def load_from_json(file_path):
    # Load the data from the JSON file
    with open(file_path, 'r') as file:
        data = json.load(file)

    return data

# Example usage
# Assuming you have lists named 'losses' and 'time' containing your data
lists_to_save = ['cumaccuracies', 'cumlosses']
save_path = '/content/my_data.json' # Replace with your desired path

save_to_json(lists_to_save, save_path)

# Load the data back from the JSON file
loaded_data = load_from_json(save_path)

# Access the loaded data
print(loaded_data['cumaccuracies']) # Access the 'losses' list
print(loaded_data['cumlosses'])    # Access the 'time' list
```

```
[]
[]
```

The training phase. Due to the problem complexity, a sample by sample approach with the gradient tape has been taken. The accuracies and the losses are saved in the loop. Every 1000 samples, the three models are saved.

The model is trained with shuffled input data.

```
In [ ]: import random
import itertools
# Get all possible pairs of indices using permutations
permutations_list = list(itertools.permutations(range(len(images)), 2))
# Randomly shuffle the list of pairs
random.shuffle(permutations_list)
print(permutations_list[:5])
```

```
[(17, 65), (51, 9), (8, 39), (86, 39), (96, 45)]
```

```
In [ ]: # Training configuration
optimizer = tf.keras.optimizers.Adam()
cumaccuracies=[]
cumlosses=[]

# Training loop
epochs = 1
k=0
for epoch in range(epochs):
    print(f"Epoch {epoch+1}/{epochs}")
    for i1, i2 in permutations_list:
        k+=1
        print("sample",k)

        x1,x2=images[i1],images[i2]
        x1 = tf.convert_to_tensor(x1)
        x1 = tf.cast(x1 , dtype=tf.float32)
        x2 = tf.convert_to_tensor(x2)
        x2 = tf.cast(x2 , dtype=tf.float32)

        augmented_images_x1 = data_augmentation([x1])
        augmented_images_x2 = data_augmentation([x2])
        augmented_images_x1 = tf.convert_to_tensor(augmented_images_x1)
        augmented_images_x1 = tf.cast(augmented_images_x1 , dtype=tf.float32)
        augmented_images_x2 = tf.convert_to_tensor(augmented_images_x2)
        augmented_images_x2 = tf.cast(augmented_images_x2 , dtype=tf.float32)
        # Randomly stack augmented images
        p = np.random.binomial(1, 0.5)
        if p == 0:
            hat_x = ([augmented_images_x1[0], augmented_images_x2[0]])

        else:
            hat_x = ([augmented_images_x2[0], augmented_images_x1[0]])

        # Calculate true order p
        true_p = np.array([p], dtype=np.float32)
        true_p = tf.convert_to_tensor(true_p)
        # Train sender
        with tf.GradientTape() as tape:
            tape.watch(sender_model.trainable_variables)
            tape.watch(emb_model.trainable_variables)
            tape.watch(rec_model.trainable_variables)
            b_x = sender_model((tf.expand_dims(x1, axis=0),tf.expand_dims(x2, axis=0))
            xi = emb_model(b_x)

            hat_p = rec_model((tf.expand_dims(hat_x[0],axis=0),tf.expand_dims(hat_x[1

            loss = (hat_p-true_p)**2
```

```

gradients = tape.gradient(loss, sender_model.trainable_variables + emb_model.
optimizer.apply_gradients(zip(gradients, sender_model.trainable_variables +
cumlosses+=[float(tf.squeeze(loss))])
predicted_p = int(tf.squeeze(hat_p) > 0.5)
true_p = int(tf.squeeze(true_p))
if predicted_p == true_p:
    cumaccuracies += [1]
else:
    cumaccuracies += [0]

if k%1000==0:
    lists_to_save = ['cumaccuracies', 'cumlosses']
    save_path = '/content/my_data.json' # Replace with your desired path

    save_to_json(lists_to_save, save_path)
    sender_model.save("send model")
    emb_model.save("emb model")
    rec_model.save("receiver model")

```

The step counter from 0 to 9900 has been erased for pdf printing.

```

In [ ]: import random
import itertools
from tensorflow.keras.models import load_model
import json
import tensorflow as tf
from tensorflow.keras import layers, models, applications
from tensorflow.keras import Model, losses
from tensorflow.keras.datasets import cifar10
import numpy as np
import matplotlib.pyplot as plt

sender_model = load_model("send model")
emb_model = load_model("emb model")
rec_model = load_model("receiver model")

# Load the data back from the JSON file
loaded_data = load_from_json(save_path)

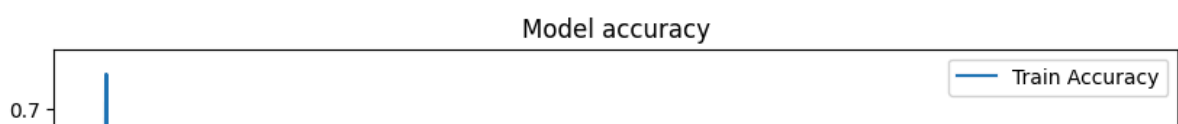
# Access the loaded data
cumaccuracies=loaded_data['cumaccuracies'] # Access the 'losses' list
cumlosses=loaded_data['cumlosses']


```

```

In [ ]: # Plot training & validation loss values
partacc=[]
for i in range(len(cumaccuracies)):
    partacc+=[sum(cumaccuracies[0:i])/(i+1)]
plt.figure(figsize=(10, 5))
plt.plot(partacc, label='Train Accuracy')
#plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='upper right')
plt.show()

```





The model is learning and was still learning at the end of the first epoch.

A test is done on 50 unseen pairs.

The test is done with a hard selection for x_1 and x_2 .

```
In [ ]: testaccuracies=[]
testlosses=[]

for i in range(50):
    x1,x2=images_test[i],images_test[i+50]
    x1 = tf.convert_to_tensor(x1)
    x1 = tf.cast(x1 , dtype=tf.float32)
    x2 = tf.convert_to_tensor(x2)
    x2 = tf.cast(x2 , dtype=tf.float32)

    augmented_images_x1 = data_augmentation([x1])
    augmented_images_x2 = data_augmentation([x2])
    augmented_images_x1 = tf.convert_to_tensor(augmented_images_x1)
    augmented_images_x1 = tf.cast(augmented_images_x1 , dtype=tf.float32)
    augmented_images_x2 = tf.convert_to_tensor(augmented_images_x2)
    augmented_images_x2 = tf.cast(augmented_images_x2 , dtype=tf.float32)
    # Randomly stack augmented images
    p = np.random.binomial(1, 0.5)
    if p == 0:
        hat_x = ([augmented_images_x1[0], augmented_images_x2[0]])
    else:
        hat_x = ([augmented_images_x2[0], augmented_images_x1[0]])

    # Calculate true order p
    true_p = np.array([p], dtype=np.float32)
    true_p = tf.convert_to_tensor(true_p)
    # Train sender
    b_x = sender_model((tf.expand_dims(x1, axis=0),tf.expand_dims(x2, axis=0)))
    b_x = int(b_x>0.5) # Hard Selection
    b_x = tf.cast(tf.expand_dims(b_x, axis=-1), dtype=tf.float32)
    xi = emb_model(b_x)
    xi = tf.expand_dims(xi, axis=0)
    hat_p = rec_model((tf.expand_dims(hat_x[0],axis=0),tf.expand_dims(hat_x[1],axis=0)),

    loss = (hat_p-true_p)**2
    testlosses+=[float(tf.squeeze(loss))]
```

```

predicted_p = int(tf.squeeze(hat_p) > 0.5)
true_p = int(tf.squeeze(true_p))
if predicted_p == true_p:
    testaccuracies += [1]
else:
    testaccuracies += [0]

print("Accuracy on the test set" , sum(testaccuracies)/len(testaccuracies))

```

Accuracy on the test set 0.64

The accuracy is above the random guess. If more time was given for training, the accuracy on the test set should be better.

```

In [ ]: import random
import itertools
from tensorflow.keras.models import load_model
import json
import tensorflow as tf
from tensorflow.keras import layers, models, applications
from tensorflow.keras import Model, losses
from tensorflow.keras.datasets import cifar10
import numpy as np
import matplotlib.pyplot as plt

sender_model = load_model("send model")
emb_model = load_model("emb model")
rec_model = load_model("receiver model")

def load_from_json(file_path):
    # Load the data from the JSON file
    with open(file_path, 'r') as file:
        data = json.load(file)

    return data

save_path = '/content/my_data.json'

# Load the data back from the JSON file
loaded_data = load_from_json(save_path)

# Access the loaded data
cumaccuracies=loaded_data['cumaccuracies'] # Access the 'losses' list
cumlosses=loaded_data['cumlosses']

```

```

In [ ]: x1l=[]
x2l=[]
aix1l=[]
aix2l=[]
bx1l=[]
pl=[]
hat_pl=[]

for i in range(10):
    x1,x2=images_test[i],images_test[i+50]
    x1l+=x1
    x2l+=x2
    x1 = tf.convert_to_tensor(x1)
    x1 = tf.cast(x1 , dtype=tf.float32)
    x2 = tf.convert_to_tensor(x2)
    x2 = tf.cast(x2 , dtype=tf.float32)

```

```

augmented_images_x1 = data_augmentation([x1])
augmented_images_x2 = data_augmentation([x2])
aix1l+=[augmented_images_x1]
aix2l+=[augmented_images_x2]

augmented_images_x1 = tf.convert_to_tensor(augmented_images_x1)
augmented_images_x1 = tf.cast(augmented_images_x1 , dtype=tf.float32)
augmented_images_x2 = tf.convert_to_tensor(augmented_images_x2)
augmented_images_x2 = tf.cast(augmented_images_x2 , dtype=tf.float32)
# Randomly stack augmented images
p = np.random.binomial(1, 0.5)
if p == 0:
    hat_x = ([augmented_images_x1[0], augmented_images_x2[0]])

else:
    hat_x = ([augmented_images_x2[0], augmented_images_x1[0]])
pl+=[p]

# Calculate true order p
true_p = np.array([p], dtype=np.float32)
true_p = tf.convert_to_tensor(true_p)
# Train sender
b_x = sender_model((tf.expand_dims(x1, axis=0),tf.expand_dims(x2, axis=0)))
b_x = int(b_x>0.5) # Hard Selection
b_x = tf.cast(tf.expand_dims(b_x, axis=-1), dtype=tf.float32)
xi = emb_model(b_x)
xi = tf.expand_dims(xi, axis=0)
hat_p = rec_model((tf.expand_dims(hat_x[0],axis=0),tf.expand_dims(hat_x[1],axis=0),
bxl+=[b_x]

xi = emb_model(b_x)
hat_pl+=[hat_p]

loss = (hat_p-true_p)**2

```

```

In [ ]: # Visualize the results
n = 5
plt.figure(figsize=(20, 8))
#image_size = 64
for i in range(n):
    # Display original
    ax = plt.subplot(7, n, i + 1)
    plt.imshow(x1l[i])
    title_str = (
        " xi : " + str(np.array(bxl[i])[0])
        + "\n p : "
        + str(np.array(pl[i]))
        + "\n  $\hat{p}$  : "
        + str(np.array(hat_pl[i][0][0][0][0]))
        + "\n \n Original x1")
    plt.title(title_str)
    plt.axis('off')

    # Display original
    ax = plt.subplot(7, n, i + 1+2*n)
    plt.imshow(x2l[i])
    plt.title("Original x2")
    plt.axis('off')

# Displav

```

```

"""
ax = plt.subplot(7, n, i + 1 + 4*n)
plt.imshow(aix1l[i][0])
plt.title("Augmented x1")
plt.axis('off')

```

```

# Display
ax = plt.subplot(7, n, i + 1 + 6*n)
plt.imshow(aix2l[i][0])
title_str = (
    "Augmented x2 : "
)
plt.title(title_str)
plt.axis('off')

```

```
plt.show()
```

xi : 0.0
p : 1
 \hat{p} : 0.58745986



xi : 1.0
p : 0
 \hat{p} : 0.47704563



xi : 0.0
p : 0
 \hat{p} : 0.61624044



xi : 0.0
p : 0
 \hat{p} : 0.5938362



xi : 1.0
p : 1
 \hat{p} : 0.29968858



Augmented x1



Augmented x1



Augmented x1



Augmented x1



Augmented x1



Augmented x2 :



Augmented x2 :



Augmented x2 :



Augmented x2 :



Augmented x2 :



To check and avoid overfitting. The different models can be saved separately during the training phase in a list every 1000 steps. The ones with the best generalization on the test set can be kept.

In []:

```

"""
#during training
if k%1000==0:
    sender_model_list+=[sender_model]
    emb_model_list+=[emb_model]
    rec_model_list+=[rec_model]

#then for the generalization tests
for i in range(len(sender_model_list):
    sender_model = sender_model_list[i]
    emb_model = emb_model_list[i]
    rec_model = rec_model_list[i]
"""

```

